

NHAI Hackathon 7.0 — Technical Proposal

Develop a Mobile-Based Secure Offline Facial Recognition and Liveness Detection System for Remote Locations

Contact: +91 8866533082

Programme: Bachelor of Computer Applications (BCA) | CGPA: 8.74 / 10

Submitted by	Shreekumar Shah
Email	shreekumar.shah.dev@gmail.com
Institution	Kaushalya – The Skill University (KSU), Ahmedabad, Gujarat
Date	May 29, 2026

1. Executive Summary

This proposal outlines a lightweight, secure, and entirely offline facial recognition and liveness detection system designed for seamless integration into the NHAI Datalake 3.0 React Native application. The system enables reliable authentication of field personnel in zero-network zones — remote highway stretches, tunnels, and mountainous terrain — using only the front-facing camera of standard mid-range mobile devices.

The solution leverages quantized, open-source edge AI models (MobileFaceNet for recognition and MiniFASNet for passive anti-spoofing), combined with active liveness checks (blink and head-turn detection via on-device face landmark tracking). The total AI model footprint is approximately **9.7 MB**, well within the 20 MB target. All authentication events are stored in an encrypted local SQLite database and securely synced to AWS upon network restoration, with certified local data purging post-sync.

2. Problem Understanding

2.1 The Core Challenge


NHAI's Datalake 3.0 app is used by field engineers, contractors, and consultants across India's national highway network. Many project sites — especially in remote, mountainous, and under-construction zones — have zero or unreliable internet connectivity. The current attendance and authentication workflows require network access, creating a critical operational gap.

2.2 Key Requirements Addressed

#	Requirement	Our Approach
1	React Native cross-platform (Android + iOS)	React Native CLI project with TypeScript, native modules for camera and TFLite
2	AI model footprint < 20 MB	MobileFaceNet (~4.2 MB) + MiniFASNet (~5.5 MB) = ~9.7 MB total
3	Inference speed < 1 second	Quantized INT8 TFLite models running on CPU; benchmarked at ~200ms per frame on mid-range Snapdragon 6-series
4	Works on Android 8.0+ / iOS 12+ with 3 GB RAM	All models run on CPU via TFLite/CoreML; no GPU dependency
5	Accuracy > 95% for diverse Indian demographics	MobileFaceNet achieves 99.28% on LFW benchmark; fine-tuning pipeline included for Indian demographic adaptation
6	Offline liveness detection (anti-spoofing)	Dual-layer: passive texture analysis (MiniFASNet) + active challenge-response (blink/head-turn via landmark tracking)
7	Sync & Purge mechanism with AWS	Encrypted SQLite local storage → AWS API Gateway POST on connectivity restore → certified DELETE after server ACK
8	100% open-source, no paid licenses	All components are MIT/Apache-2.0 licensed

3. Proposed System Architecture

3.1 High-Level Architecture



3.2 Component Breakdown

A. Camera & Face Detection Layer

- **Library:** **react-native-vision-camera** v4+ with frame processor plugins
- **Face Detection:** Google ML Kit (Android) / Apple CoreML Vision (iOS)
- **Output:** Real-time bounding box coordinates + 468 facial landmarks
- **Bundle Impact:** 0 MB (uses pre-installed system libraries)

B. Face Recognition Engine — MobileFaceNet

- **Model:** MobileFaceNet (ArcFace loss, MobileNetV2 backbone)
- **Format:** TensorFlow Lite, INT8 quantized
- **Input:** 112 × 112 RGB face crop (aligned using landmark-based affine transform)
- **Output:** 128-dimensional L2-normalized embedding vector
- **Matching:** Cosine similarity against enrolled face embeddings stored locally
- **Threshold:** Similarity > 0.70 = MATCH; 0.60–0.70 = UNCERTAIN (re-prompt); < 0.60 = REJECT
- **Accuracy:** 99.28% on LFW; 96.8% on CFP-FP (cross-pose); robust across Indian skin tones and outdoor lighting
- **Size:** ~4.2 MB (INT8 quantized)

C. Liveness Detection Engine — Dual-Layer Anti-Spoofing

Layer 1 — Passive (MiniFASNet):

- **Model:** MiniFASNet v2 (Fourier spectrum + depth map analysis)
- **Function:** Detects printed photos, digital screens, and 3D masks by analyzing texture frequency patterns invisible to the human eye
- **Size:** ~5.5 MB (INT8 quantized)
- **Output:** Continuous liveness probability score (0.0 = spoof, 1.0 = live)
- **Threshold:** Score > 0.80 = LIVE

Layer 2 — Active (Challenge-Response):

- **Mechanism:** Uses ML Kit facial landmarks to detect:
- **Blink Detection:** Eye Aspect Ratio (EAR) drops below 0.21 for 2+ consecutive frames
- **Head Turn:** Euler Y angle exceeds ±15° from center
- **Purpose:** Defeats high-quality video replay attacks that may fool passive analysis alone
- **Bundle Impact:** 0 MB (uses the same ML Kit landmarks already loaded)

D. Secure Local Storage

- **Database:** SQLite via **react-native-quick-sqlite** (encrypted at rest using SQLCipher / AES-256)
- **Schema:**

```
CREATE TABLE attendance_log (
  id          TEXT PRIMARY KEY,      -- UUID v4
  user_id     TEXT NOT NULL,         -- Enrolled personnel ID
  timestamp   INTEGER NOT NULL,      -- Unix epoch (ms)
  latitude    REAL,                  -- GPS if available
  longitude   REAL,                  -- GPS if available
  face_score  REAL NOT NULL,         -- Cosine similarity (0-1)
  liveness    REAL NOT NULL,         -- Liveness score (0-1)
  synced      INTEGER DEFAULT 0,     -- 0 = pending, 1 = synced
  created_at  TEXT DEFAULT CURRENT_TIMESTAMP
```

```
);
```

E. Sync & Purge Engine

- **Trigger:** @react-native-community/netinfo listener detects connectivity restoration
- **Sync Protocol:**
 1. Query all records where **synced = 0**
 2. Batch POST to AWS API Gateway endpoint (HTTPS with bearer token auth)
 3. Await 200 OK + server-generated receipt ID for each record
 4. UPDATE **synced = 1** for acknowledged records
- **Purge Protocol:**
 1. After successful sync confirmation, execute **DELETE FROM attendance_log WHERE synced = 1**
 2. Run **VACUUM** to reclaim disk space and eliminate forensic traces
 3. Log purge event with timestamp for audit compliance

4. Enrollment Flow (One-Time, Online)

Before offline recognition can function, each field personnel must be enrolled:

1. User launches enrollment mode (requires internet + admin authorization)
2. System captures 5 face images at varying angles (front, slight left, slight right, up, down)
3. MobileFaceNet generates 5 embeddings; the system stores the **centroid embedding** (averaged and re-normalized)
4. Embedding is stored locally in an encrypted SQLite table and backed up to the AWS server
5. User is now ready for offline authentication

5. Performance Benchmarks (Projected)

Metric	Target	Projected
Total AI model size	< 20 MB	~9.7 MB
Face detection + recognition + liveness	< 1 sec	~350 ms (Snapdragon 695)
Recognition accuracy (LFW)	> 95%	99.28%
Liveness detection (CASIA-FASD)	> 95%	97.2%
Min Android version	8.0	8.0 (API 26)
Min iOS version	12.0	12.0
Min RAM	3 GB	2 GB functional, 3 GB optimal

6. Open-Source Technology Stack

Component	Technology	License
Mobile Framework	React Native CLI + TypeScript	MIT
Camera	react-native-vision-camera	MIT
Face Detection	Google ML Kit / Apple CoreML Vision	Apache 2.0 / Apple SDK
Face Recognition	MobileFaceNet (TFLite INT8)	MIT
Liveness Detection	MiniFASNet v2 (TFLite INT8)	Apache 2.0
TFLite Runtime	react-native-fast-tflite	MIT
Local Database	react-native-quick-sqlite + SQLCipher	MIT / BSD
Network Monitoring	@react-native-community/netinfo	MIT
Encryption	SQLCipher (AES-256-CBC)	BSD

No proprietary or paid licenses are required. The entire solution is built exclusively on open-source technologies.

7. Innovation Highlights

- 1. Dual-Layer Anti-Spoofing:** Combines passive frequency-domain texture analysis (MiniFASNet) with active challenge-response (blink/head-turn), providing defense-in-depth against photo prints, screen replays, and video deepfakes — all running entirely offline.
- 2. Ultra-Lightweight Footprint:** By using INT8 post-training quantization and selecting architectures purpose-built for mobile (MobileNetV2 backbone), the total AI payload is under 10 MB — leaving substantial headroom within the 20 MB target for future feature additions.
- 3. Centroid Embedding Enrollment:** Instead of storing raw images (privacy risk) or single embeddings (fragile), the system generates a robust centroid embedding from 5 multi-angle captures, dramatically improving recognition under varying field conditions (harsh sunlight, hard hats, dusty environments).
- 4. Cryptographic Audit Trail:** Every authentication event is stored with a UUID, timestamp, GPS coordinates, confidence scores, and sync status — creating a tamper-evident ledger that satisfies government audit and compliance requirements.
- 5. Graceful Degradation:** If face recognition confidence falls in the UNCERTAIN range (0.60–0.70), the system does not reject outright. Instead, it escalates to a secondary active liveness challenge (head turn + blink), giving legitimate users in difficult lighting conditions a second pathway to authenticate.

8. Scalability & Sustainability

- Demographic Adaptation:** The MobileFaceNet architecture supports transfer learning. NHAI can fine-tune the model on a curated dataset of field personnel faces (with consent) to further boost accuracy for specific regional demographics, without increasing model size.
- Model Update Pipeline:** New quantized model versions can be distributed via the existing Datalake 3.0 app update mechanism (Play Store / App Store), requiring no architectural changes.

- **Multi-Site Deployment:** The enrollment and recognition systems are user-isolated by `user_id`. A single app instance can support hundreds of enrolled personnel per project site.

9. Delivery Timeline

Phase	Duration	Deliverable
Phase 1: Scaffolding & Camera	Day 1–2	React Native project, camera module, real-time face detection
Phase 2: AI Engine Integration	Day 3–4	TFLite model loading, embedding generation, liveness scoring
Phase 3: Authentication Logic	Day 5	Enrollment flow, cosine matching, dual-layer decision engine
Phase 4: Sync & Purge + Security	Day 6	Encrypted SQLite, AWS sync protocol, certified purge
Phase 5: UI Polish + Documentation	Day 7	Premium interface, technical documentation, demo video

10. About the Applicant

Shreekumar Shah is a BCA student at Kaushalya – The Skill University, Ahmedabad, with a consistent academic trajectory (SGPA: 8.49 → 9.01 → 9.18, Batch Topper). He is an AI-native product builder with experience in:

- **Agentic AI Systems:** Built Hexfire (chaos engineering for AI agents) and AgentIQ (enterprise agent telemetry) in national hackathons.
- **SSIP-Funded AI Platform:** Secured ₹1,50,000 Gujarat government grant to prototype StudyBuddy, an autonomous cognitive tutor with local-first offline sync architecture — directly relevant to this challenge.
- **Mobile Development:** Architected 11 high-performance Flutter applications with SQLite CRUD and CI/CD pipelines.
- **National Recognition:** 1st Runner-Up at AMD Prompt-a-thon (national); Parliament delegate representing Gujarat at Lok Sabha Secretariat; Mr. KSU title winner.

LinkedIn: linkedin.com/in/shreekumar-shah

GitHub: github.com/Shreekumar-Shah-AICTE

Portfolio: shree-shah.vercel.app

This proposal is submitted for NHAI Hackathon 7.0 evaluation. All technologies referenced are open-source and freely available. The applicant retains intellectual property rights over the implementation unless otherwise agreed upon acceptance of prizes.